
Heuristic Search (Informed Search)

Heuristic Search

- In **uninformed search**, we don't try to evaluate which of the nodes on the frontier are most promising. We never “look-ahead” to the goal.

E.g., in uniform cost search we always expand the cheapest path. We don't consider the cost of getting to the goal.
- Often we have some other knowledge about the merit of nodes, e.g., going the wrong direction in Romania.

Heuristic Search

Merit of a frontier node: different notions of merit.

- If we are concerned about the **cost of the solution**, we might want a notion of merit of how costly it is to get to the goal from that search node.
- If we are concerned about **minimizing computation** in search we might want a notion of ease in finding the goal from that search node.
- We will focus on the “**cost of solution**” notion of merit.

Heuristic Search

- The idea is to develop a domain specific heuristic function $h(n)$.
- $h(n)$ guesses the cost of getting to the goal from node n .
- There are different ways of guessing this cost in different domains. I.e., heuristics are domain specific.

Heuristic Search

- Convention: If $h(n_1) < h(n_2)$ this means that we guess that it is cheaper to get to the goal from n_1 than from n_2 .
- We require that
 - $h(n) = 0$ for every node n that satisfies the goal.
 - Zero cost of getting to a goal node from a goal node.

Using only $h(n)$: Greedy best-first search

- We use $h(n)$ to rank the nodes on the frontier.
 - Always expand node with lowest h -value.
- We are greedily trying to achieve a low cost solution.
- However, this method **ignores the cost of getting to n** , so it can be lead astray exploring nodes that cost a lot to get to but seem to be close to the goal:


→ step cost = 10

→ step cost = 100

$h(n1) = 70$

$h(n3) = 50$

(Incomplete)



Using only $h(n)$: Greedy best-first search.

- We use $h(n)$ to rank the nodes on the frontier.
 - Always expand node with lowest h -value.
- We are greedily trying to achieve a low cost solution.
- However, this method **ignores the cost of getting to n** , so it can be lead astray exploring nodes that cost a lot to get to but seem to be close to the goal:


→ step cost = 10

→ step cost = 100

$h(n1) = 70$

$h(n3) = 50$

(Incomplete)



A* search

- Take into account the cost of getting to the node as well as our estimate of the cost of getting to the goal from n .
- Define an evaluation function $f(n)$
 $f(n) = g(n) + h(n)$
 - $g(n)$ is the cost of the path to node n
 - $h(n)$ is the heuristic estimate of the cost of getting to a goal node from n .
- Always expand the node with lowest f -value on the frontier.
- The f -value is an estimate of the cost of getting to the goal via this node (path).

Conditions on $h(n)$

- We want to analyze the behavior of the resultant search.
- Completeness, time and space, optimality?
- To obtain such results we must put some further conditions on the heuristic function $h(n)$ and the search space.

Conditions on $h(n)$: Admissible

- We always assume that $c(n_1 \rightarrow n_2) \geq \epsilon > 0$. The cost of any transition is greater than zero and can't be arbitrarily small.
- Let $h^*(n)$ be the **cost of an optimal path** from n to a goal node (∞ if there is no path). Then an **admissible** heuristic satisfies the condition
$$h(n) \leq h^*(n)$$
i.e. h always underestimates of the true cost.
- Hence
 - $h(g) = 0$, for any goal node, g
 - $h(n) = \infty$ if there is not path from n to a goal node

Consistency (aka monotonicity)

- Is a stronger condition than $h(n) \leq h^*(n)$.
- A **monotone/consistent** heuristic satisfies the triangle inequality (for all nodes n_1, n_2):
$$h(n_1) \leq c(n_1 \rightarrow n_2) + h(n_2)$$
- Note that there might be more than one transition (action) between n_1 and n_2 , the inequality must hold for all of them.
- Note that monotonicity implies admissibility. Why?

Intuition behind admissibility

$h(n) \leq h^*(n)$ means that the search won't miss any promising paths.

- If it really is cheap to get to a goal via n (i.e., both $g(n)$ and $h^*(n)$ are low), then $f(n) = g(n) + h(n)$ will also be low, and the search won't ignore n in favor of more expensive options.
- This can be formalized to show that admissibility implies optimality.

Intuition behind monotonicity

$$h(n1) \leq c(n1 \rightarrow n2) + h(n2)$$

- This says something similar, but in addition one won't be “locally” mislead. See next example.

Example: admissible but nonmonotonic

The following h is not consistent since $h(n2) > c(n2 \rightarrow n4) + h(n4)$.
But it is admissible.

→ step cost = 200
→ step cost = 100

$$g(n) + h(n) = f(n)$$


$\{S\} \rightarrow \{n1 [200+50=250], n2 [200+100=300]\}$
 $\rightarrow \{n2 [100+200=300], n3 [400+50=450]\}$
 $\rightarrow \{n4 [200+50=250], n3 [400+50=450]\}$
 $\rightarrow \{goal [300+0=300], n3 [400+50=450]\}$

$h(n1) = 50$

$h(n3) = 50$

$h(n2) = 200$

$h(n4) = 50$



We **do find** the optimal path as the heuristic is still admissible. **But** we are mislead into ignoring n2 until after we expand n1.

Example: admissible but nonmonotonic

The following h is not consistent since $h(n2) > c(n2 \rightarrow n4) + h(n4)$.
But it is admissible.

→ step cost = 200
→ step cost = 100

$$g(n) + h(n) = f(n)$$


$\{S\} \rightarrow \{n1 [200+50=250], n2 [200+100=300]\}$
 $\rightarrow \{n2 [100+200=300], n3 [400+50=450]\}$
 $\rightarrow \{n4 [200+50=250], n3 [400+50=450]\}$
 $\rightarrow \{goal [300+0=300], n3 [400+50=450]\}$

$h(n1) = 50$

$h(n3) = 50$

$h(n2) = 200$

$h(n4) = 50$



We **do find** the optimal path as the heuristic is still admissible. **But** we are mislead into ignoring n2 until after we expand n1.

Consequences of monotonicity

1. The f-values of nodes along a path must be non-decreasing.

- Let $\langle \text{Start} \rightarrow n_1 \rightarrow n_2 \dots \rightarrow n_k \rangle$ be a path. We claim that
$$f(n_i) \leq f(n_{i+1})$$

- Proof:

$$\begin{aligned} f(n_i) &= c(\text{Start} \rightarrow \dots \rightarrow n_i) + h(n_i) \\ &\leq c(\text{Start} \rightarrow \dots \rightarrow n_i) + c(n_i \rightarrow n_{i+1}) + h(n_{i+1}) \\ &= c(\text{Start} \rightarrow \dots \rightarrow n_i \rightarrow n_{i+1}) + h(n_{i+1}) \\ &= g(n_{i+1}) + h(n_{i+1}) \\ &= f(n_{i+1}). \end{aligned}$$

Consequences of monotonicity

2. If n_2 is expanded after n_1 , then $f(n_1) \leq f(n_2)$
(the f -value increases monotonically)

Proof:

- If n_2 was on the frontier when n_1 was expanded, then $f(n_1) \leq f(n_2)$ otherwise we would have expanded n_2 .
- If n_2 was added to the frontier after n_1 's expansion, then let n be an ancestor of n_2 that was present when n_1 was being expanded (this could be n_1 itself). We have $f(n_1) \leq f(n)$ since A^* chose n_1 while n was present in the frontier. Also, since n is along the path to n_2 , by property (1) we have $f(n) \leq f(n_2)$. So, we have $f(n_1) \leq f(n_2)$.

Consequences of monotonicity

3. When n is expanded every path with lower f -value has already been expanded.

- **Proof:** Assume by contradiction that there exists a path $\langle \text{Start}, n_0, n_1, n_{i-1}, n_i, n_{i+1}, \dots, n_k \rangle$ with $f(n_k) < f(n)$ and n_i is its last expanded node.
 - n_{i+1} must be on the frontier while n is expanded, so
 - a) by (1) $f(n_{i+1}) \leq f(n_k)$ since they lie along the same path.
 - b) since $f(n_k) < f(n)$ so we have $f(n_{i+1}) < f(n)$
 - c) by (2) $f(n) \leq f(n_{i+1})$ because n is expanded before n_{i+1} .
 - Contradiction from b&c!

Consequences of monotonicity

4. With a monotone heuristic, the first time A^* expands a state, it has found the minimum cost path to that state.

Proof:

- Let **PATH1** = **<Start, n0, n1, ..., nk, n>** be **the first** path to n found. We have $f(\text{path1}) = c(\text{PATH1}) + h(n)$.
- Let **PATH2** = **<Start, m0, m1, ..., mj, n>** be another path to n found later. we have $f(\text{path2}) = c(\text{PATH2}) + h(n)$.
- By property (3), $f(\text{path1}) \leq f(\text{path2})$
- hence: $c(\text{PATH1}) \leq c(\text{PATH2})$

Consequences of monotonicity

Complete.

- Yes, consider a least cost path to a goal node
 - $\text{SolutionPath} = \langle \text{Start} \rightarrow n_1 \rightarrow \dots \rightarrow G \rangle$ with cost $c(\text{SolutionPath})$
 - Since each action has a cost $\geq \epsilon > 0$, there are only a finite number of paths that have cost $\leq c(\text{SolutionPath})$.
 - All of these paths must be explored before any path of cost $> c(\text{SolutionPath})$.
 - So eventually SolutionPath , or some equal cost path to a goal must be expanded.

Time and Space complexity.

- When $h(n) = 0$, for all n h is monotone.
 - A^* becomes uniform-cost search!
- It can be shown that when $h(n) > 0$ for some n , the number of nodes expanded can be no larger than uniform-cost.
- Hence the same bounds as uniform-cost apply. (These are worst case bounds). Still exponential unless we have a very very good h !
- In real world problems, we run out of time and memory! IDA^* used to address memory issues.

Consequences of monotonicity

Optimality


- Yes, by (4) the first path to a goal node must be optimal.

Cycle Checking

- If we do cycle checking (multiple path checking) e.g. using **GraphSearch** instead of **TreeSearch**, it is still optimal. Because by property (4) we need keep only the first path to a node, rejecting all subsequent paths.

Search generated by monotonicity

Gradually adds “ f -contours” of nodes (cf. breadth-first adds layers)
Inside each counter, the f values are less than or equal to counter value!



- For uniform cost search, bands are “circular”.
- With more accurate heuristics, bands stretch out more toward the goal.

Admissibility without monotonicity

When “h” is admissible but not monotonic.

- Time and Space complexity remain the same. Completeness holds.
- Optimality still holds (without cycle checking), but need a different argument: don't know that paths are explored in order of cost.

Proof of optimality (without cycle checking):

- Assume the goal path $\langle S, \dots, G \rangle$ found by A^* has cost bigger than the optimal cost: i.e. $C^*(G) < f(G)$.
- There must exist a node n in the optimal path that is still in the frontier.
- We have: $f(n) = g(n) + h(n) \leq g(n) + h^*(n) = C^*(G) < f(G)$
why?
- Therefore, $f(n)$ must have been selected **before** G by A^* . contradiction!

Admissibility without monotonicity

What about Cycle Checking?

- No longer guaranteed we have found an optimal path to a node *the first time* we visit it.
- So, cycle checking might not preserve optimality.
 - To fix this: for previously visited nodes, must remember cost of previous path. If new path is cheaper must explore again.
- contours of monotonic heuristics don't hold.

Space Problems with A*

- A* has the same potential space problems as BFS
- IDA* - Iterative Deepening A* is similar to Iterative Lengthening Search and similarly addresses space issues.

IDA* - Iterative Deepening A*

Objective: reduce memory requirements for A*

- Like iterative deepening, but now the “cutoff” is the f-value ($g+h$) rather than the depth
- At each iteration, the cutoff value is the smallest f-value of any node that exceeded the cutoff on the previous iteration
- Avoids overhead associated with keeping a sorted queue of nodes
- Two new parameters:
 - curBound (any node with a bigger f-value is discarded)
 - smallestNotExplored (the smallest f-value for discarded nodes in a round) when frontier becomes empty, the search starts a new round with this bound.

Building Heuristics: Relaxed Problem

- One useful technique is to consider an easier problem, and let $h(n)$ be the cost of reaching the goal in the easier problem.
- 8-Puzzle moves.
 - Can move a tile from square A to B if
 - A is adjacent (left, right, above, below) to B
 - **and** B is blank
- Can relax some of these conditions
 1. can move from A to B if A is adjacent to B (ignore whether or not position is blank)
 2. can move from A to B if B is blank (ignore adjacency)
 3. can move from A to B (ignore both conditions).

Building Heuristics: Relaxed Problem

- **#3** *“can move from A to B (ignore both conditions)”*.

leads to the **misplaced tiles** heuristic.

- To solve the puzzle, we need to move each tile into its final position.
- Number of moves = number of misplaced tiles.
- Clearly $h(n)$ = number of misplaced tiles \leq the $h^*(n)$ the cost of an optimal sequence of moves from n .

- **#1** *“can move from A to B if A is adjacent to B (ignore whether or not position is blank)”*

leads to the **manhattan distance** heuristic.

- To solve the puzzle we need to slide each tile into its final position.
- We can move vertically or horizontally.
- Number of moves = sum over all of the tiles of the number of vertical and horizontal slides we need to move that tile into place.
- Again $h(n)$ = sum of the manhattan distances $\leq h^*(n)$
 - in a real solution we need to move each tile at least that far and we can only move one tile at a time.

Building Heuristics: Relaxed Problem

The optimal cost to nodes in the relaxed problem is an **admissible heuristic** for the original problem!

Proof: the optimal solution in the original problem is a (*not necessarily optimal*) solution for relaxed problem, therefore it must be at least as expensive as the optimal solution in the relaxed problem.

Building Heuristics: Relaxed Problem

Comparison of IDS and A* (average total nodes expanded):

Depth	IDS	A*(Misplaced) h1	A*(Manhattan) h2
10	47,127	93	39
14	3,473,941	539	113
24	---	39,135	1,641

Let **h1**=Misplaced, **h2**=Manhattan

- Does **h2 always** expand fewer nodes than **h1**?
 - Yes! Note that **h2 dominates h1**, i.e. for all n : $h1(n) \leq h2(n)$. From this you can prove **h2** is faster than **h1**.
 - Therefore, among several admissible heuristic the one with highest value is the fastest.

Building Heuristics: Pattern databases.

- Admissible heuristics can also be derived from solution to **subproblems**: Each state is mapped into a partial specification, e.g. in 15-puzzle only *position of specific tiles matters*.

- Here are goals for two subproblems (called Corner and Fringe) of 15puzzle.

- Note** the location of BLANK!




Fig. 2. The Fringe and Corner Target Patterns.

- By searching **backwards** from these goal states, we can compute the distance of any configuration of these tiles to their goal locations. We are ignoring the identity of the other tiles.
- For any state n , the number of moves required to get these tiles into place form a lower bound on the cost of getting to the goal from n .

Building Heuristics: Pattern databases.

- These configurations are stored in a database, along with the number of moves required to move the tiles into place.
- The **maximum** number of moves taken **over all of the databases** can be used as a heuristic.
- On the 15-puzzle
 - The fringe data base yields about a 345 fold decrease in the search tree size.
 - The corner data base yields about 437 fold decrease.
- Sometimes **disjoint patterns** can be found, then the number of moves can be **added** rather than taking the max (if we only count moves of the target tiles).

Local Search

- So far, we keep the paths to the goal.
- For some problems (like 8-queens) we don't care about the path, we **only care about the solution**. Many real problem like Scheduling, IC design, and network optimizations are of this form.
- **Local search** algorithms operate using a single **current state** and generally move to neighbors of that state.
- There is an **objective function** that tells the value of each state. The goal has the highest value (global maximum).
- Algorithms like **Hill Climbing** try to move to a neighbour with the highest value.
- Danger of being stuck in a **local maximum**. So some randomness is added to “shake” out of local maxima.
- **Simulated Annealing**: Instead of the best move, take a random move and if it improves the situation then always accept, otherwise accept with a probability <1 .
- [If interested read these two algorithms from the R&N book].